

Hands-on Lab: Monitoring and Optimizing Your Databases in PostgreSQL

Estimated time needed: 45 minutes

In this lab, you'll learn how to monitor and optimize your database in PostgreSQL with both the command line interface (CLI) and database administration tool, pgAdmin.

Objectives

After completing this lab, you will be able to:

1. Monitor the performance of your database with the command line interface and pgAdmin.
2. Identify optimal data types for your database.
3. Optimize your database via the command line with best practices.

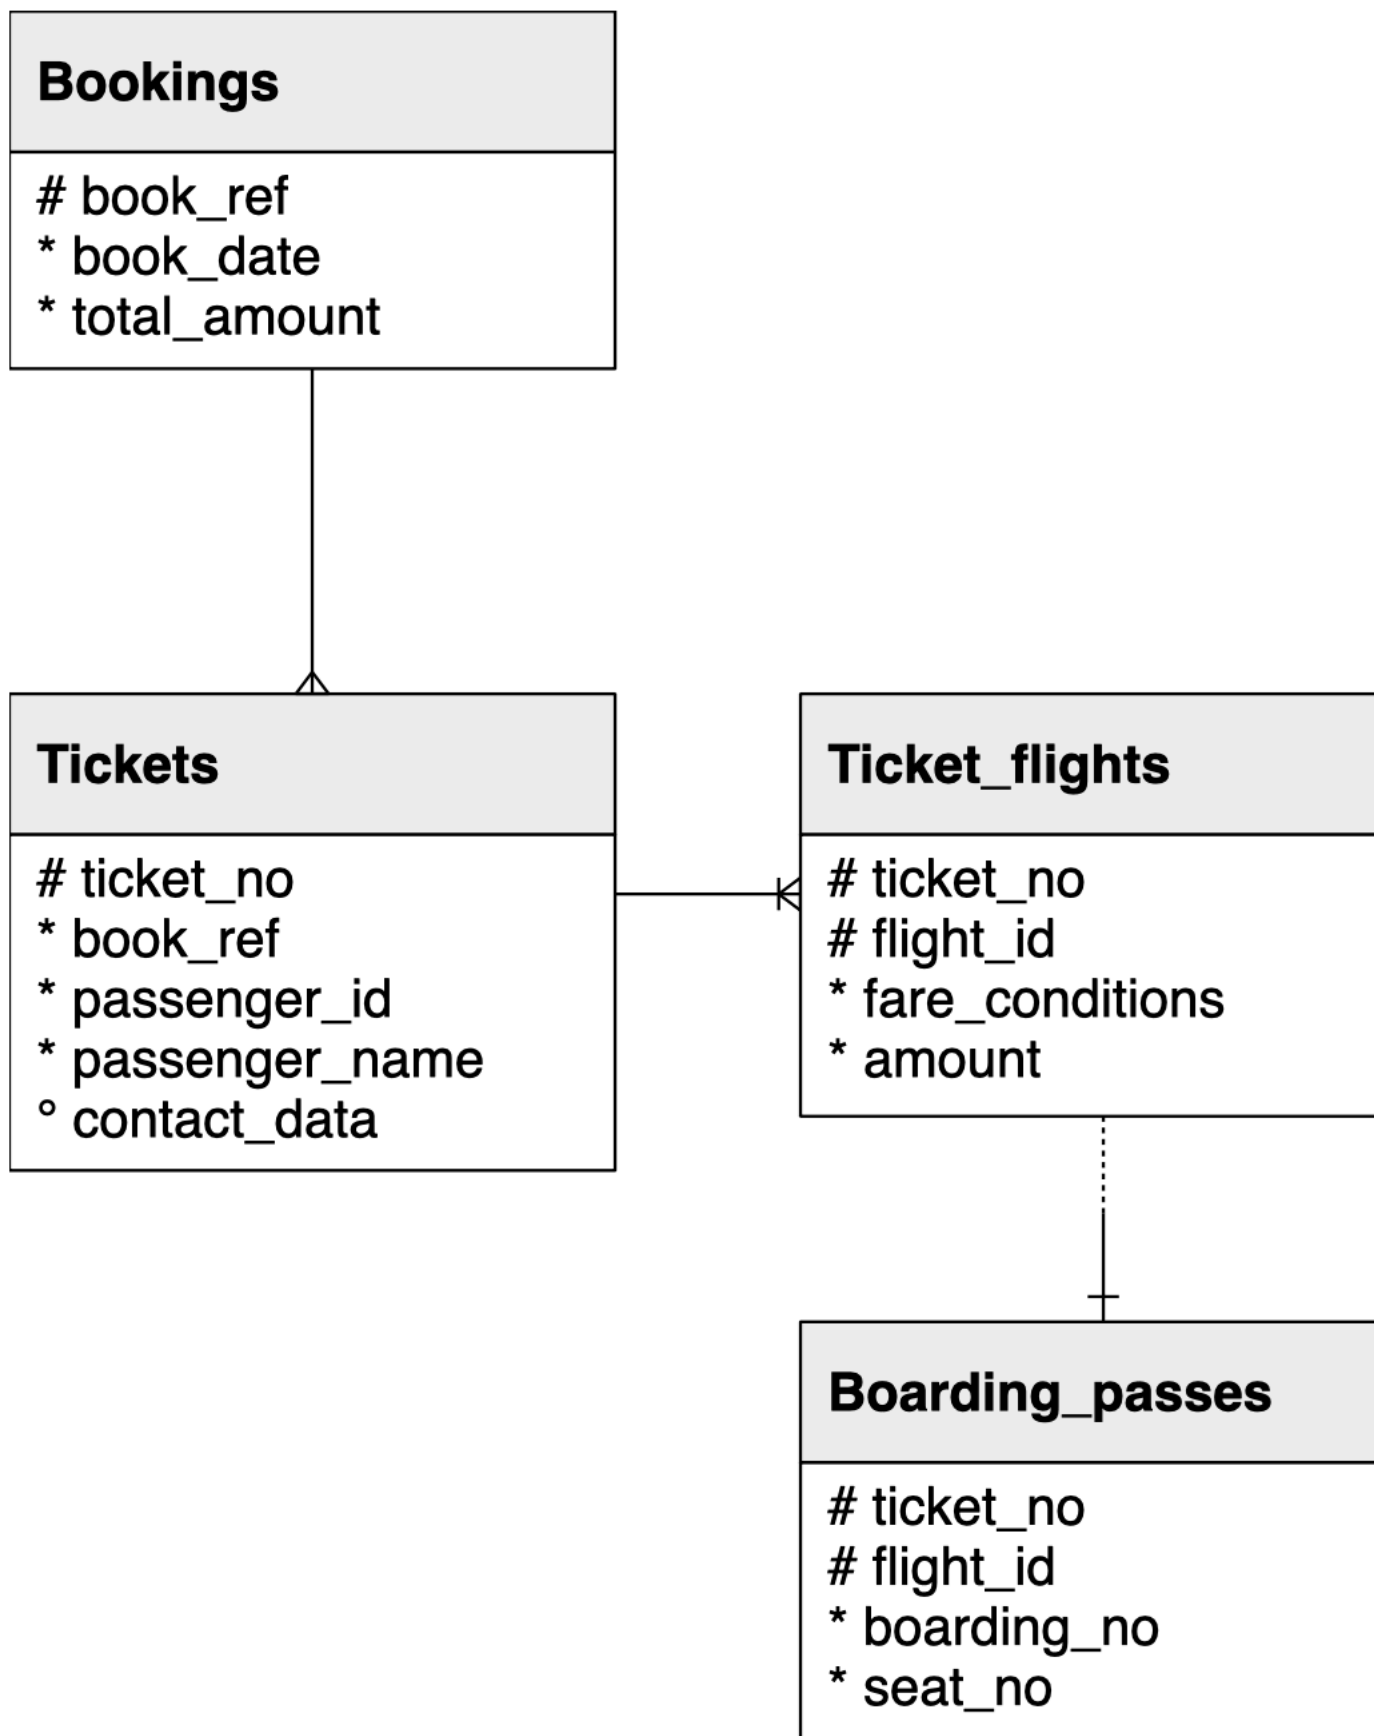
Software Used in this Lab

In this lab, you will be using PostgreSQL. It is a popular open source object relational database management system (RDBMS) capable of performing a wealth of database administration tasks, such as storing, manipulating, retrieving, and archiving data.

To complete this lab, you will be accessing the PostgreSQL service through the IBM Skills Network (SN) Cloud IDE, which is a virtual development environment you will use throughout this course.

Database Used in this Lab

In this lab, you will use a database from <https://postgrespro.com/education/demodb> distributed under the [PostgreSQL licence](#). It stores a month of data about airline flights in Russia and is organized according to the following schema:



Exercise 1: Create Your Database

To get started with this lab, you'll launch PostgreSQL in Cloud IDE and create our database with the help of a SQL file.

Task A: Start PostgreSQL in Cloud IDE

1. To start PostgreSQL, navigate to the **Skills Network Toolbox**, select **Databases**, and select **PostgreSQL**.

Select **Start**. This will start a session of PostgreSQL in Skills Network Labs.




The screenshot shows the Skills Network Labs interface. The top menu bar includes File, Edit, Selection, View, Go, Run, Terminal, and Help. The sidebar on the left contains icons for various tools, with the 'Skills Network Toolbox' icon highlighted by a red box and labeled '1'. The main panel displays a list of databases under the 'DATASETS' section, with 'PostgreSQL INACTIVE' selected by a red box and labeled '2'. The right panel shows the PostgreSQL configuration page, with the 'Start' button highlighted by a red box and labeled '4'. The 'Start' button is blue with white text. The right panel also shows the PostgreSQL version (v13.2) and the connection information (v5.0).

The **Inactive** label will change to **Starting**. This may take a minute or so.

2. When the label changes to **Active**, it means your session has started.

PostgreSQL

ACTIVE

 v13.2 |  v5.0 |  v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI




New Terminal

Task B: Create Your Database

1. Open a new terminal by selecting the **New Terminal** button near the bottom of the PostgreSQL tab.

PostgreSQL

ACTIVE

 v13.2 |  v5.0 |  v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details


Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin



Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

2. With the terminal, you'll want to download the **demo** database that you're using in this lab. This database contains a month of data about flights in Russia.

To download it, you can use the following command:

```
wget https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/example-guided-project/flights_RUSSIA_small.sql
```

You should now see the SQL file in your file explorer in Cloud IDE.

The screenshot shows the Visual Studio Code interface. The Explorer sidebar on the left displays the 'EXPLORER: PROJECT' view with a folder named 'postgres' containing a file 'flights_RUSSIA_small.sql', which is highlighted with a red rectangle. The main editor area shows the 'PostgreSQL' extension interface with a 'Stop' button and sections for 'Summary' and 'Connection Information'. The terminal window at the bottom shows a command line session where the file 'flights_RUSSIA_small.sql' is being processed, including a message about resolving 'cf-courses-data.s3.amazonaws.com' and saving the file.

3. Let's return to the PostgreSQL tab and select the **PostgreSQL CLI** button near the bottom of the tab. This button will open a new PostgreSQL command line session.

PostgreSQL

ACTIVE

v13.2

v5.0

v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

4. Now, you want to import the data from the file that you downloaded.

- ▶ [Hint \(Click Here\)](#)
- ▶ [Solution \(Click Here\)](#)

5. With our created database, let's see what tables you have. How many tables are there?

- ▶ [Hint \(Click Here\)](#)
- ▶ [Solution \(Click Here\)](#)

Great! With your environment and database set up, let's take a look at how you can monitor and optimize this database!

Exercise 2: Monitor Your Database

Database monitoring refers to reviewing the operational status of your database and maintaining its health and performance. With proper and proactive monitoring, databases will be able to maintain a consistent performance. Any problems that emerge, such as sudden outages, can be identified and resolved in a timely manner.

Tools such as pgAdmin, an open source graphical user interface (GUI) tool for PostgreSQL, come with several features that can help monitor your database. The main focus in this lab will be using the command line interface to monitor your database, but we'll also take a quick look at how the monitoring process can be replicated in pgAdmin.

Monitoring these statistics can be helpful in understanding your server and its databases, detecting any anomalies and problems that may arise.

Task A: Monitor Current Activity

To start, let's take a look at how you can monitor current server and database activity in PostgreSQL.

Server Activity

You can take a look at the server activity by running the following query:

```
SELECT pid, username, datname, state, state_change FROM pg_stat_activity;
```

This query will retrieve the following:

| Column | Description |
|--------------|--|
| pid | Process ID |
| username | Name of user logged in |
| datname | Name of database |
| state | Current state, with two common values being: active (executing a query) and idle (waiting for new command) |
| state_change | Time when the state was last changed |

This information comes from the **pg_stat_activity**, one of the built-in statistics provided by PostgreSQL.

1. Copy the query and paste it into the terminal.

You should see the following output:

```
demo=# SELECT pid, username, datname, state, state_change FROM pg_stat_activity;
 pid | username | datname | state | state_change
-----+-----+-----+-----+-----
  42 |          |         | idle |
  44 | postgres |         | idle |
  51 | postgres | postgres | idle | 2021-10-13 22:11:20.330154+00
1090 | postgres | demo    | active | 2021-10-13 22:11:20.725355+00
  40 |          |         | idle |
  39 |          |         | idle |
  41 |          |         | idle |
(7 rows)
```

As you can see, there are currently 7 active connections to the server, with two of them being connected to databases that you're familiar with. After all, you started in the default **postgres** database, which is now idle, and now you're actively querying in the **demo** database.

2. To see what other columns are available for viewing, feel free to take a look at the [pg_stat_activity documentation](#)!

Let's say you wanted to see all the aforementioned columns, in addition to the actual text of the query that was last executed. Which column should you add to review that?

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

Please note, if your table looks strange or squished, you can resize the terminal window by dragging it out.

If your result shows the text **(END)**, then type in q to exit that view. Whenever you encounter this view, you can use q to return to your original view.

3. With queries, you can apply filtering. What if you only wanted to see the states that were **active**? How would you do that?

- [Hint \(Click Here\)](#)
- [Solution \(Click Here\)](#)

Database Activity

When looking at database activity, you can use the following query:

```
SELECT datname, tup_inserted, tup_updated, tup_deleted FROM pg_stat_database;
```

This query will retrieve the following:

| Column | Description |
|--------------|---|
| datname | Name of database |
| tup_inserted | Number of rows inserted by queries in this database |
| tup_updated | Number of rows updated by queries in this database |
| tup_deleted | Number of rows deleted by queries in this database |

This information comes from the **pg_stat_database**, one of the statistics provided by PostgreSQL.

1. Copy the query and paste it into the terminal.

You should see the following output:


```
demo=# SELECT datname, tup_inserted, tup_updated, tup_deleted FROM pg_stat_database
 datname | tup_inserted | tup_updated | tup_deleted
-----+-----+-----+-----
 postgres |          2 |          1 |          0
 demo     |         2290162 |          22 |          0
 template1 |          0 |          0 |          0
 template0 |          0 |          0 |          0
(5 rows)
```

As you can see, the two databases that are returned are the **postgres** and **demo**. These are databases that you are familiar with.

The other two, **template1** and **template0** are default templates for databases, and can be overlooked in this analysis.

Based on this output, you now know that **demo** had about 2,290,162 rows inserted and 22 rows updated.

2. To see what other columns are available for viewing, you can read through the [pg_stat_database documentation](#).

Let's say you wanted to see the number of rows fetched and returned by this database.

Note: The number of rows fetched is the number of rows that were returned. The number of rows returned is the number of rows that were read and scanned by the query.

What query should you use to do that?

▼ Hint (Click Here)

Take a look at the documentation provided earlier. Which column(s) would show you what you need?

▼ Solution (Click Here)

To see the number of rows fetched and returned, you use the following query:

```
SELECT datname, tup_fetched, tup_returned FROM pg_stat_database;
```

Your result should look similar to the following:

```
demo=# SELECT datname, tup_fetched, tup_returned FROM pg_stat_database
 datname | tup_fetched | tup_returned
-----+-----+-----
 postgres |        8513 |        65666
 demo     |       76915 |       2554350
 template1 |          0 |          0
 template0 |          0 |          0
(5 rows)
```

Notice how the rows returned tend to be greater than the rows fetched. If you consider how tables are read, this makes sense because not all the rows scanned may be the ones that are returned.

3. With queries, you can apply filtering. What if you only wanted to see the database details (rows inserted, updated, deleted, returned and fetched) for **demo**?

▼ Hint (Click Here)

Recall that you can filter queries with the **WHERE** clause.

▼ Solution (Click Here)

To filter the results so that only those from the **demo** database are shown, you use the following query:

```
SELECT datname, tup_inserted, tup_updated, tup_deleted, tup_fetched, tup_returned FROM pg_stat_database WHERE datname = 'demo';
```

Your result should look similar to the following:

```
demo=# SELECT datname, tup_inserted, tup_updated, tup_deleted, tup_fetched, tup_returned FROM pg_stat_database WHERE datname = 'demo';
 datname | tup_inserted | tup_updated | tup_deleted | tup_fetched | tup_returned
-----+-----+-----+-----+-----+-----
 demo    |       2290162 |          22 |          0 |       588014 |       2554350
(1 row)
```

Later, we'll take a look at how you can monitor these activities in pgAdmin.

Task B: Monitor with pgAdmin

Another method of monitoring your database comes in the form of pgAdmin. In order to use this tool, you'll need to first launch it.

1. Open the PostgreSQL tab from the Skills Network Toolbox and select the pop-out button next to the **pgAdmin** button. This will open pgAdmin in a new tab.

The screenshot shows the PostgreSQL Skills Network Labs interface. At the top, it says "PostgreSQL" in large white text, followed by a green "ACTIVE" button. Below this, there are three icons representing different versions: a database icon for "v13.2", a gear icon for "v5.0", and a terminal icon for "v13.2". A message states: "Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment." Below this is a blue "Stop" button. There are three tabs: "Summary" (selected), "Connection Information", and "Details". The "Summary" tab content says: "Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section." Below this are two input fields: "Username:" and "Password:", both with blacked-out text and copy icons. A message says: "You can manage PostgreSQL via:". Below this is a blue "pgAdmin" button and a blue button with a white square and an upward-pointing arrow, which is highlighted with a red box. Below this is a message: "Or to interact with the database in the terminal, select one of these options:". At the bottom are two blue buttons: "PostgreSQL CLI" and "New Terminal".

2. In the left panel, select the dropdown next to **Servers**. You'll be prompted to enter a password.


pgAdmin File Object Tools Help

Browser

Servers (1)

1 > postgres

Welcome


 pgAdmin
Management

Feature rich | Maximizing productivity

pgAdmin is an Open Source administration tool designed to answer the needs of PostgreSQL users.

Quick Links




Getting Started

 PostgreSQL Documentation

3. Return to your Cloud IDE session. In the PostgreSQL tab, select the copy button next to the **Password** field. This is the password that you can enter into pgAdmin.

PostgreSQL

ACTIVE

 v13.2 |  v5.0 |  v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.


Stop


Summary

Connection Information

Details


Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username: 

Password: 

You can manage PostgreSQL via:

pgAdmin



Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

Paste that password into pgAdmin. Then, click **OK**.

Your server will now load.

- On the home page, under **Dashboard**, you will have access to server or database statistics, depending on which you are looking at.

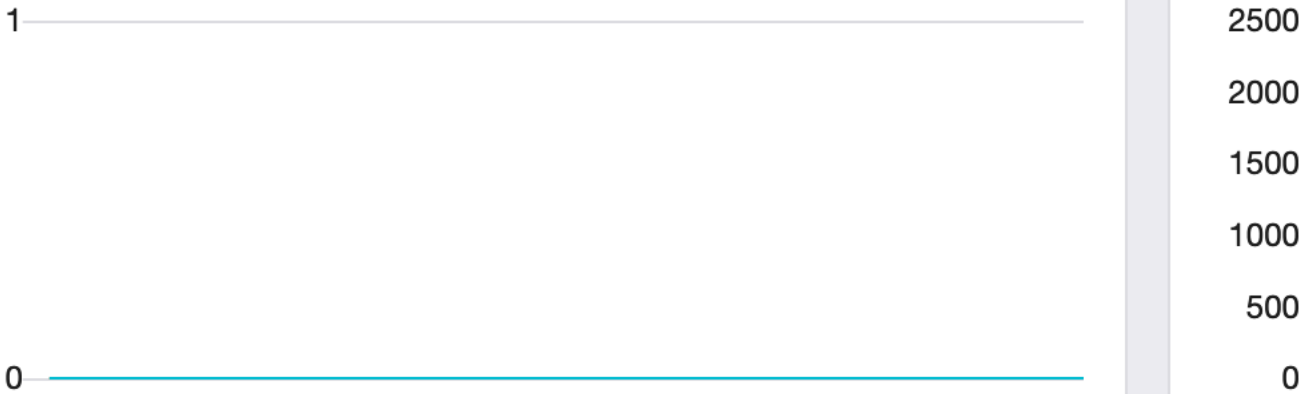
Server sessions



Tuples in

Inserts Updates Delete

Tuples



Server activity

Sessions

Locks

Prepared Transactions

Configuration

| | | | PID | Database | User | Application |
|---|---|---|-----|----------|----------|-------------------------|
| ✖ | ■ | ▶ | 40 | | | |
| ✖ | ■ | ▶ | 41 | | | |
| ✖ | ■ | ▶ | 42 | | | |
| ✖ | ■ | ▶ | 43 | | | |
| ✖ | ■ | ▶ | 45 | | postgres | |
| ✖ | ■ | ▶ | 52 | postgres | postgres | sysdig-agent |
| ✖ | ■ | ▶ | 271 | postgres | postgres | pgAdmin 4 - DB:postgres |

The table below lists the displayed statistics on the Dashboard that correspond with the statistics that you accessed with the CLI:

| Chart | Description |
|--------------------------|---|
| Server/Database sessions | Displays the total sessions that are running. For servers, this is similar to the pg_stat_activity , and for databases, this is similar to the pg_stat_database . |
| Transactions per second | Displays the commits, rollbacks, and transactions taking place. |
| Tuples in | Displays the number of tuples (rows) that have been inserted, updated, and deleted, similar to the tup_inserted , tup_updated , and tup_deleted columns from pg_stat_database . |
| Tuples out | Displays the number of tuples (rows) that have been fetched (returned as output) or returned (read or scanned). This is similar to tup_fetched and tup_returned from pg_stat_database . |
| Server activity | Displays the sessions, locks, prepared transactions, and configuration for the server. In the Sessions tab, it offers a look at the breakdown of the sessions that are currently active on the server, similar to the view provided by pg_stat_activity . To check for any new processes, you can select the refresh button at the top-right corner. |

5. You can test these charts out by starting another session.

Return to the tab with the Cloud IDE environment. On the PostgreSQL tab, select **PostgreSQL CLI**. This will start a new session of PostgreSQL with the CLI.

PostgreSQL

ACTIVE

v13.2

v5.0

v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

6. Once you have started that instance, switch back to the tab with pgAdmin.

What do you notice?

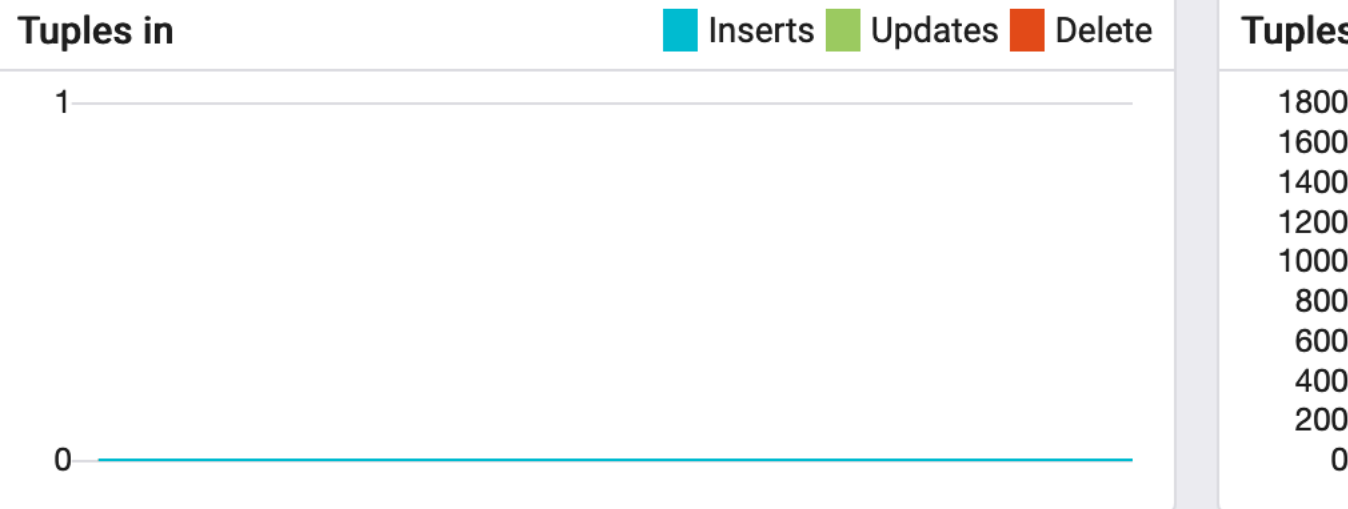

▼ Hint (Click Here)

Consider this: Which chart(s) monitors active sessions? Remember that one of the charts may need to be refreshed before updates are shown.

▼ Solution (Click Here)

You may have noticed that the **Server sessions** saw an increase of sessions. It increased from 7 to 8 sessions. This makes sense since you started a new session with PostgreSQL CLI.

To see that change reflected in **Server Activity**, you'll have to click the refresh button to see that an additional **postgres** database session appeared.



Server activity

| | | | PID | Database | User | Application |
|---|---|---|-----|----------|----------|-------------------------|
| ✖ | ■ | ▶ | 40 | | | |
| ✖ | ■ | ▶ | 41 | | | |
| ✖ | ■ | ▶ | 42 | | | |
| ✖ | ■ | ▶ | 43 | | | |
| ✖ | ■ | ▶ | 45 | | postgres | |
| ✖ | ■ | ▶ | 52 | postgres | postgres | sysdig-agent |
| ✖ | ■ | ▶ | 271 | postgres | postgres | pgAdmin 4 - DB:postgres |
| ✖ | ■ | ▶ | 415 | postgres | postgres | psql |

7. To see the dashboard for the **demoss** database, navigate to the left panel and select the **Databases** dropdown and then select the **demo** database to connect to it.

As you can see, similar statistics are displayed for the database.

pgAdmin

File ▾Object ▾Tools ▾Help ▾

Browser

Database icons

Search icon

Dashboard

Properties

SQL

1 ▾ Servers (1)

2 ▾ postgres

1 ▾ Databases (2)

2 ▾ demo

> Casts

> Catalogs

> Event Triggers

> Extensions

> Foreign Data Wrappers

> Languages

> Publications

> Schemas

> Subscriptions

> postgres

> Login/Group Roles

> Tablespaces

Database sessions

1

0

Tuples in

1

0

Server activity

Sessions

Locks

Prepared

| | | | PID | User |
|---|---|---|-----|----------|
| ✖ | ■ | ▶ | 951 | postgres |

8. Let's run a query on the database! To do that, navigate to the menu bar and select Tools > Query Tool.

You can run any query. To keep things simple, let's run the following to select all the data from the **bookings** table:

```
SELECT * FROM bookings;
```

Select the run button. You will see that the query has successfully loaded.

Dashboard

Properties

SQL

Statistics

Dependencies

Dependents

No

demo/postgres@postgres

Query Editor

Query History

1

SELECT * FROM bookings;

Data Output

Explain

Messages

Notifications

| | <div>book_ref</div> <div>[PK] character (6)</div> | <div>book_date</div> <div>timestamp with time zone</div> | <div>total_amount</div> <div>numeric (10,2)</div> |
|----|---|--|---|
| 1 | 00000F | 2017-07-05 00:12:00+00 | 265700.0 |
| 2 | 000012 | 2017-07-14 06:02:00+00 | 37900.0 |
| 3 | 000068 | 2017-08-15 11:27:00+00 | 18100.0 |
| 4 | 000181 | 2017-08-10 10:28:00+00 | 131800.0 |
| 5 | 0002D8 | 2017-08-07 18:40:00+00 | 23600.0 |
| 6 | 0002DB | 2017-07-29 03:30:00+00 | 101500.0 |
| 7 | 0002E0 | 2017-07-11 13:09:00+00 | 89600.0 |
| 8 | 0002F3 | 2017-07-10 02:31:00+00 | 69600.0 |
| 9 | 00034E | 2017-08-04 13:52:00+00 | 73300.0 |
| 10 | 000352 | 2017-07-05 23:02:00+00 | 109500.0 |
| 11 | 000374 | 2017-08-12 07:13:00+00 | 136200.0 |

9. In pgAdmin, switch back to the database's **Dashboard** tab. You can refresh the **Server activity** and check to see if any of the charts have shown a spike since the data was retrieved.

What do you notice?

▼ Hint (Click Here)

Recall what you queried. Which chart would reflect those changes?

▼ Solution (Click Here)

You may have noticed that the number of tuples (rows) returned (read/scanned) was greater than 250,000.

Database sessions

2

1

0

Tuples in

Inserts Updates Delete

1

0

Tuples

30000

25000

20000

15000

10000

5000

Server activity

Sessions

Locks

Prepared Transactions

| | | | PID | User | Application |
|---|---|---|------|----------|--------------------------|
| ✖ | ■ | ▶ | 951 | postgres | pgAdmin 4 - DB:demo |
| ✖ | ■ | ▶ | 1496 | postgres | pgAdmin 4 - CONN:5716674 |

You can check the number of rows scanned with EXPLAIN:

```
EXPLAIN SELECT * FROM bookings;
```

If you can't see the full text, simply drag the **QUERY PLAN** column out.

This statement reveals that 262,788 rows were scanned, which is similar to the amount that was read/scanned based on the spike in the **Tuples out** chart.

Dashboard Properties SQL Statistics Dependencies Dependents

demo/postgres@postgres

Query Editor Query History

```
1 EXPLAIN SELECT * FROM bookings;
```

Data Output Explain Messages Notifications

| QUERY PLAN | | |
|------------|--|--|
| | text | |
| 1 | Seq Scan on bookings (cost=0.00..4301.88 rows=262788 width=21) | |

While you can monitor your database through the command line alone, tools like pgAdmin can be helpful in providing a visual representation of how your server and its databases are performing.

PostgreSQL also offers logging capabilities to monitor and troubleshoot your lab, which will be further discussed in the Troubleshooting lab.

Note: This Task cannot be executed in this lab environment and can be done only locally

Monitor Performance Over Time

Extensions, which can enhance your PostgreSQL experience, can be helpful in monitoring your database. One such extension is **pg_stat_statements**, which gives you an aggregated view of query statistics.

1. To enable the extension, enter the following command:

```
CREATE EXTENSION pg_stat_statements;
```

This will enable the **pg_stat_statements** extension, which will start to track the statistics for your database.

2. Now, let's edit the PostgreSQL configuration file to include the extension you just added:

```
ALTER SYSTEM SET shared_preload_libraries = 'pg_stat_statements';
```

For the changes to take effect, you will have to restart your database. You can do that by typing `exit` in the terminal to stop your current session. Close the terminal and return to the PostgreSQL tab. Select **Stop**.

PostgreSQL

ACTIVE

v13.2 | v5.0 | v13.2

Connect to PostgreSQL and pgAdmin directly in your Skills Network Labs environment.

Stop

Summary

Connection Information

Details

Your database and pgAdmin server are now ready to use and available with the following PostgreSQL, please check out the Details section.

Username:

Password:

You can manage PostgreSQL via:

pgAdmin

Or to interact with the database in the terminal, select one of these options:

PostgreSQL CLI

New Terminal

When the session has become **Inactive** once more, select **Start** to restart your session.

3. Once your session has started, open the **PostgreSQL CLI**.

You'll need to reconnect to the **demo** database, which you can do by using the following command:

```
\connect demo
```

4. You can see if this extension has been loaded by checking both the installed extensions and the **shared_preload_libraries**.

First let's check the installed extensions:

```
\dx
```

```
demo=# \dx
```

| Name | Version | Schema | List of installed extensions | Description |
|--------------------|---------|------------|--|-------------|
| pg_stat_statements | 1.8 | bookings | track planning and execution statistics of all SQL statements executed | |
| plpgsql | 1.0 | pg_catalog | PL/pgSQL procedural language | |
| (2 rows) | | | | |

Notice how **pg_stat_statements** has been installed.

You can also check the **shared_preload_libraries** with:

```
show shared_preload_libraries;
```

```
demo=# show shared_preload_libraries;
shared_preload_libraries
-----
pg_stat_statements
(1 row)
```

pg_stat_statements is also shown under **shared_preload_libraries**.

5. Since the results returned by **pg_stat_statements** can be quite long, let's turn on expanded table formatting with the following command:

```
\x
```

This will display the output tables in an expanded table format.

```
demo=# \x
Expanded display is on.
demo=#
```

You can turn it off by repeating the `\x` command.

6. From the [pg_stat_statements documentation](#), you'll see the various columns available to be retrieved.

Let's say you wanted to retrieve the database ID, the query, and total time that it took to execute the statement (in milliseconds).

▼ Hint (Click Here)

Recall how you selected columns previously to display statistics.

▼ Solution (Click Here)

Use the following query to extract the details you'd like to retrieve:

```
SELECT dbid, query, total_exec_time FROM pg_stat_statements;
```

You'll notice that you can scroll through the results, which may look similar to the following:

```
-[ RECORD 10 ]--
dbid          | 13442
query         | COMMIT
total_exec_time | 0.7378020000000001
-[ RECORD 11 ]--
dbid          | 16384
query         | SELECT e.extname AS "Name", e.extversion AS "Version", n.nspname AS "Schema", c.description AS "Description"
+
+
+
FROM pg_catalog.pg_extension e LEFT JOIN pg_catalog.pg_namespace n ON n.oid = e.extnamespace LEFT JOIN pg_catalog.pg_class c ON c.oid = e.extclass AND c.classoid = $1::pg_catalog.regclass
+
+
ORDER BY 1
total_exec_time | 0.32195999999999997
-[ RECORD 12 ]--
dbid          | 16384
query         | show shared_preload_libraries
total_exec_time | 0.015136
(END)
```

Unlike **pg_stat_activity**, which showed the latest query that was run, **pg_stat_statements** shows an aggregated view of the queries that were run since the extension was installed.

7. What if you wanted to check which database name matches the database ID?

▼ Hint (Click Here)

Consider how you previously retrieved information about the database.

▼ Solution (Click Here)

Use the following query to extract the database ID and database name:

```
SELECT oid, datname FROM pg_database;
```

```
demo=# SELECT oid, datname FROM pg_database
demo=# ;
-[ RECORD 1 ]-----
oid          | 13442
datname      | postgres
-[ RECORD 2 ]-----
oid          | 16384
datname      | demo
-[ RECORD 3 ]-----
oid          | 1
datname      | template1
-[ RECORD 4 ]-----
oid          | 13441
datname      | template0
```

Based on this, you can now see that database ID **16384** is the **demo** database. This makes sense because you performed the query `show shared_preload_libraries` on the **demo** database, which appeared in **pg_stat_statements**.

It's important to note that adding these extensions can increase your server load, which may affect performance. If you need to drop the extension, you can achieve that with the following command:


```
DROP EXTENSION pg_stat_statements;
```

If you check the current extensions with `\dx`, you'll also see that **pg_stat_statements** no longer appears.

You should also reset the `shared_preload_libraries` in the configuration file:

```
ALTER SYSTEM RESET shared_preload_libraries;
```

After this, you'll need to exit the terminal and restart the PostgreSQL CLI to see the changes reflected in `show shared_preload_libraries;`.

Exercise 3: Optimize Your Database

Data optimization is the maximization of the speed and efficiency of retrieving data from your database. Optimizing your database will improve its performance, whether that's inserting or retrieving data from your database. Doing this will improve the experience of anyone interacting with the database.

Similar to MySQL, there are optimal data types and maintenance (otherwise known as "vacuuming") that can be applied to optimize databases.

Task A: Optimize Data Types

When it comes to optimizing data types, understanding the data values will help in selecting the proper data type for the column.

Let's take a look at an example in the **demo** database.

- 1. Return to the CLI session that you opened previously (or open a new session if it has been closed).

If you're no longer connected to the **demo** database, you can reconnect to it!

▼ Hint (Click Here)

Remember how you connected to the database earlier in this lab.

▼ Solution (Click Here)

You can use the following command to connect to the **demo** database:

```
\connect demo
```

- 2. Let's list out the tables in the database with the following command:

```
\dt
```

demo=# \dt

| List of relations | | | |
|-------------------|-----------------|-------|----------|
| Schema | Name | Type | Owner |
| bookings | aircrafts_data | table | postgres |
| bookings | airports_data | table | postgres |
| bookings | boarding_passes | table | postgres |
| bookings | bookings | table | postgres |
| bookings | flights | table | postgres |
| bookings | seats | table | postgres |
| bookings | ticket_flights | table | postgres |
| bookings | tickets | table | postgres |
| (8 rows) | | | |

- 3. Now that you know which tables are in the database, select the first one, **aircrafts_data** and see what data you can pull from it. How can you select all of its data?

▼ Hint (Click Here)

Which statement can you use to select all the data in this table?

▼ Solution (Click Here)

You can use the following query to select all the data from **aircrafts_data**:

```
SELECT * FROM aircrafts_data;
```

```
demo=# SELECT * FROM aircrafts_data;
 aircraft_code |          model          | range
-----+-----+-----
 773           | {"en": "Boeing 777-300"} | 11100
 763           | {"en": "Boeing 767-300"} |  7900
 SU9           | {"en": "Sukhoi Superjet-100"} |  3000
 320           | {"en": "Airbus A320-200"} |  5700
 321           | {"en": "Airbus A321-200"} |  5600
 319           | {"en": "Airbus A319-100"} |  6700
 733           | {"en": "Boeing 737-300"} |  4200
 CN1           | {"en": "Cessna 208 Caravan"} |  1200
 CR2           | {"en": "Bombardier CRJ-200"} |  2700
(9 rows)
```

You can see that there are 9 entries in total with three columns: **aircraft_code**, **model**, and **range**.

4. For the purposes of this lab, we'll create a hypothetical situation that will potentially require changing the data types of columns to optimize them.

Let's say that **aircraft_code** is always set to three characters, **model** will always be in a JSON format and **range** has a maximum value of 12,000 and minimum value of 1,000.

In this case, what would be the best data types for each column?

▼ Hint (Click Here)

Take a look at the [Data Types Documentation](#) by PostgreSQL to see which data types would fit the columns!

▼ Solution (Click Here)

Based on the documentation, the following data types would be suitable for the following columns:

1. **aircraft_code**: char(3), since you know that the aircraft code will always be fixed to three characters.
2. **model**: json, which is a special data type that PostgreSQL supports.
3. **range**: smallint, since the range of its numbers falls between -32,768 to 32,767.

•

You can check the current data types (and additional details such as the indexes and constraints) of the **aircrafts_data** table with the following:

```
\d aircrafts_data
```

```
demo=# \d aircrafts_data;
Table "bookings.aircrafts_data"
  Column      |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 aircraft_code | character(3)    |           | not null |
 model         | jsonb           |           | not null |
 range         | integer         |           | not null |
Indexes:
    "aircrafts_pkey" PRIMARY KEY, btree (aircraft_code)
Check constraints:
    "aircrafts_range_check" CHECK (range > 0)
Referenced by:
    TABLE "flights" CONSTRAINT "flights_aircraft_code_fkey" FOREIGN KEY
    TABLE "seats" CONSTRAINT "seats_aircraft_code_fkey" FOREIGN KEY (ai
```

•

Notice that most of the columns in this table have been optimized for our sample scenario, except for the **range**. This may be because the range was unknown in the original database.

For this lab, let's take the opportunity to optimize that column for your hypothetical situation. You can do this by changing the data type of the column.

Please note that in this lab you'll first need to drop a view, which is another way our data can be presented, in order to change the column's data type. Otherwise, you will encounter an error. This is a special case for this database because you loaded a SQL file that included commands to create views. In your own database, you may not need to drop a view.

To drop the **aircrafts** view, use the following command:

```
DROP VIEW aircrafts;
```

To change the column's data type, you'll use the following command:

```
ALTER TABLE aircrafts_data ALTER COLUMN range TYPE smallint;
```

aircrafts_data is the table you want to change and **range** is the column you want to change to data type **smallint**.

Now, let's check the table's columns and data types again!

► Hint (Click Here)

▼ Solution (Click Here)

With the following command, you can check the columns and data types of the **aircrafts_data** table:

```
\d aircrafts_data
```

```
demo=# \d aircrafts_data;
Table "bookings.aircrafts_data"
  Column      |      Type      | Collation | Nullable | Default
-----+-----+-----+-----+-----
 aircraft_code | character(3)    |           | not null |
 model         | jsonb           |           | not null |
 range         | smallint        |           | not null |
Indexes:
    "aircrafts_pkey" PRIMARY KEY, btree (aircraft_code)
Check constraints:
    "aircrafts_range_check" CHECK (range > 0)
Referenced by:
    TABLE "flights" CONSTRAINT "flights_aircraft_code_fkey" FOREIGN KEY (aircraft_code) REFERENCES bookings.aircrafts_data (aircraft_code)
    TABLE "seats" CONSTRAINT "seats_aircraft_code_fkey" FOREIGN KEY (aircraft_code) REFERENCES bookings.aircrafts_data (aircraft_code)
```

You can see that the data type has successfully been changed, optimizing your table in this hypothetical situation.

Task B: Vacuum Your Databases

In your day-to-day life, you can vacuum our rooms to keep them neat and tidy. You can do the same with databases by maintaining and optimizing them with some vacuuming.

In PostgreSQL, vacuuming means to clean out your databases by reclaiming any storage from “dead tuples”, otherwise known as rows that have been deleted but have not been cleaned out.

Generally, the **autovacuum** feature is automatically enabled, meaning that PostgreSQL will automate the vacuum maintenance process for you.

You can check if this is enabled with the following command:

```
show autovacuum;
```

```
demo=# show autovacuum;
autovacuum
-----
on
(1 row)
```

As you can see, **autovacuum** is enabled.

Since **autovacuum** is enabled, let's check to see when your database was last vacuumed.

To do that, you can use the **pg_stat_user_tables**, which displays statistics about each table that is a user table (instead of a system table) in the database. The columns that are returned are the same ones listed in [pg_stat_all_tables documentation](#).

What if you wanted to check the table (by name), the estimated number of dead rows that it has, the last time it was autovacuumed, and how many times it has been autovacuumed?

▼ Hint (Click Here)

Recall how you can select specific columns from statistics.

▼ Solution (Click Here)

To select the table name, number of dead rows, the last time it was autovacuumed, and the number of times this table has been autovacuumed, you can use the following query:

```
SELECT relname, n_dead_tup, last_autoanalyze, autovacuum_count FROM pg_stat_user_tables;
```

```
demo=# SELECT relname, n_dead_tup, last_autoanalyze, autovacuum_count FROM pg_stat_user_tables;
 relname      | n_dead_tup |      last_autoanalyze      | autovacuum_count |
-----+-----+-----+-----+
 tickets      |          0 | 2021-10-14 15:14:22.586399+00 |          1 |
 boarding_passes |          0 | 2021-10-14 15:14:19.709957+00 |          1 |
 seats        |          0 | 2021-10-14 15:14:19.967282+00 |          1 |
 aircrafts_data |          0 |                               |          0 |
 flights      |          0 | 2021-10-14 15:14:19.926375+00 |          1 |
 ticket_flights |          0 | 2021-10-14 15:14:16.945501+00 |          1 |
 bookings     |          0 | 2021-10-14 15:14:22.842128+00 |          1 |
 airports_data |          0 | 2021-10-14 15:14:19.946403+00 |          0 |
(8 rows)
```

Notice that you currently don’t have any “dead tuples” (deleted rows that haven’t yet been cleaned out) and so far, these tables have been autovacuumed once. This makes sense given that the database was just created and based on the logs, autovacuumed then.

Conclusion

Congratulations! Now, not only do you know how to monitor and optimize your database with the CLI, but you can also do so with pgAdmin. You will now be able to apply this knowledge to any PostgreSQL databases you create and modify in the future.

Author(s)

Kathy An

Other Contributor(s)

Changelog

| Date | Version | Changed by | Change Description |
|------------|---------|---------------|-------------------------|
| 2021-10-14 | 1.0 | Kathy An | Created initial version |
| 2023-05-04 | 1.1 | Rahul Jaideep | Updated Markdown file |

© IBM Corporation 2023. All rights reserved.